

AI & ML 101

For AEs

Scott Reed

July 2025

© 2022 Snowflake Inc. All Rights Reserved 1 Agenda

Part 1: AI & ML Fundamentals

Part 2: Snowflake's View & Where to

Compete

© 2022 Snowflake Inc. All Rights Reserved

2 Part 1: AI & ML Fundamentals

© 2022 Snowflake Inc. All Rights Reserved Brief History of Machine Learning

Pre-2000: The Foundation Years Post-2000: The Data Explosion

- 1950s: Birth of Machine Learning
 - Term "machine learning" coined (IBM)
 - "Turing Test" to measure machine intelligence
- 1960s-1970s: The AI Winter
 - Early chatbots and pattern recognition (ELIZA)
 - Computers not strong enough
- 1980s-1990s: The Comeback
 - Neural network training breakthrough
 - IBM Deep Blue beats chess champion (1997)

Smart ideas, but waiting for technology to catch

- 2000s: The Data Explosion
 - Proliferation of massive datasets
 - Computing power catches up
- 2010s: Deep Learning Dominance
 - AI beats humans at complex games
 - ML becomes part of everyday software
- 2020s: Generative AI Era
 - ChatGPT brings AI to mainstream
 - Anyone can create with AI tools
 - AI as common as smartphones

Technology caught up, data exploded, AI went
mainstream
up

© 2022 Snowflake Inc. All Rights Reserved

4Let's clear this up...

Citation

© 2022 Snowflake Inc. All Rights Reserved 5Comparison to Similar Disciplines

Dimension Economics Statistics Machine Learning

Primary Goal

Explain & predict human/market

behavior to guide policy and

strategy

Draw reliable inferences about

populations from samples

Maximize predictive accuracy and

automate decisions

Theory-driven hypotheses (e.g.,

Hypothesis testing framework

Starting Point

supply & demand models)

with probability models

Data-driven pattern discovery

with minimal assumptions

Data & Tools

Published economic indicators;

Stata & MATLAB

Representative samples /

studies; R, SAS

Large, often messy,

high-dimensional datasets;

Python, ML Platforms, Notebooks

Key Output

Causal insights and policy

recommendations

Confidence intervals, p-values,

significance tests

Real-time predictions,

classification scores,

recommendations

Interpretability vs.

Accuracy

High emphasis on transparent

causal stories

High interpretability; statistical

significance required

Trades interpretability for higher

accuracy in complex models

© 2022 Snowflake Inc. All Rights Reserved 6Two Main Personas

Data Scientist / Research Scientist • Frame business problems and conduct

exploratory data analysis (EDA)

- Build prototype models and validate

statistical insights

- Communicate findings through visualizations

presentations

- Handoff validated prototypes to engineering

teams

- Tools: SQL, python, advanced math, ML

platforms, notebooks, autoML

© 2022 Snowflake Inc. All Rights Reserved

ML Engineer / ML Ops Engineer

- Refactor prototypes into scalable,

production-ready pipelines (CI/CD)

- Optimize model training, inference cost,

and latency performance

- Build APIs, containerize services, and

implement monitoring systems

- Handle real-time drift detection and

automated rollback systems

- Tools: Python, Git, Docker / Kubernetes,

Infrastructure as code

Glossary of Key Terms

Term Definition

Algorithm

The mathematical "recipe" that tells a computer how to learn patterns from data

Model

The object that is created once an algorithm learns from data—what makes actual predictions. This is a static object.

Model Training

Teaching a model using historical data so it can make accurate predictions on new data

Model Scoring /

Inference

When the trained model makes predictions on new, real-world data—where business value happens

Accuracy

How often the model makes correct predictions, usually expressed as a percentage

Feature

Individual data points or columns the model uses to make decisions (age, purchase history, etc.)

© 2022 Snowflake Inc. All Rights Reserved 8

Glossary of Key Terms (continued)

Term Definition

Feature

Engineering

Selecting and creating the most useful data inputs—often what separates good from great models

Label/Target The correct answer you want the model to predict ("will customer churn?", "optimal price?")

Hyperparameters Settings that control how the algorithm learns—like temperature and time when baking

Cross-Validation Testing technique to ensure the model will work on new data, not just training data

AutoML Technology that automates model-building, making ML accessible to non-experts

Drift Performance degradation over time—either data changes or accuracy drops, requiring retraining

A/B Testing Comparing two model versions to determine which performs better in real business conditions

© 2022 Snowflake Inc. All Rights Reserved 9Types of Machine Learning Techniques

Citation

© 2022 Snowflake Inc. All Rights Reserved 10Make it make

sense...

Supervised learning builds off

labels and the ultimate goal is

to build a model to predict

which “class” a new record

would fall into

Unsupervised learning takes

raw data without labels and

finds patterns in the data to

form logical “groupings”

Citation

© 2022 Snowflake Inc. All Rights Reserved 11Use Case to ML Technique Mapping

SHOPPER 360

OMNI PLANNING &

FORECASTING

MODERNIZING THE

RETAIL SUPPLY CHAIN

STORE OPERATIONAL

INTELLIGENCE

How to design a

customer centricity

How to forecast product

demand, optimize

How to design supply

chain visibility, improve

How to surface intraday

sales and operational

strategy?

pricing, and reduce

risk?

accuracy and prioritize

immediate actions?

performance metrics to
help support the store
teams understand
current sales patterns.

Customer Segmentation:

Unsupervised

Customer Lifetime Value

(CLV):

Supervised

Customer Churn:

Supervised

Demand Forecasting:

Supervised

Dynamic Pricing:

Supervised

Fraud Detection:

Supervised

RETAIL & CPG

DATA CLOUD

© 2025 Snowflake Inc. All Rights Reserved

Inventory Optimization /

Demand Forecasting:

Supervised

Anomaly Detection in

Logistics:

Unsupervised

Shipment ETA:

Supervised

Market Basket

Analysis:

Unsupervised

Foot Traffic & Sales

Forecasting:

Supervised

RETAIL MEDIA

How to monetize data

through digital

advertising and shopper

insights?

Ad Targeting &

Click-Through Rate:

Supervised

Audience Segmentation:

Unsupervised

Personalized

Promotions:

SupervisedPart 2: Snowflake's View &

Where to Compete

© 2022 Snowflake Inc. All Rights ReservedModel Lifecycle in Snowflake

UNIFIED GOVERNANCE

Develop

& Iterate

Orchestrate

Monitor

& Automate

Manage Deploy

& Serve

© 2025 Snowflake Inc. All Rights ReservedDevelop

& Iterate

Orchestrate

& Automate

Manage Deploy

& Serve

Monitor

What is it? Why it Matters Discovery Questions

During this phase, Data Scientists

are doing EDA, feature

engineering, model building and

testing.

The bulk of ML use cases spend a

good amount of time in this stage.

This is a critical group to win over

and there is high compute
upside with daily, consistent
usage via notebooks, model
training, and feature engineering
pipelines.

Snowflake notebooks with
container runtime, distributed
APIs, along with git integration
makes us very competitive

© 2025 Snowflake Inc. All Rights Reserved

How much time is spent moving,
sampling, or waiting for data versus
actually building and iterating on
models?

How do you ensure fresh,
production-quality data without
creating security risks or data
governance challenges?

When using open-source libraries
like Scikit-learn or PyTorch, how do
you manage the underlying
compute environments and
package dependencies? Develop
& Iterate

Orchestrate

& Automate

Manage Deploy

& Serve

Monitor

What is it? Why it Matters Discovery Questions

In this phase, ML engineers are

taking code developed by data

scientists and turning it into

repeatable, production-grade

pipelines.

This involves running multi-step

data preparation and model

training workflows through

Snowpark and ML Jobs

(container runtime).

Simplifying the “hand-off” between

data scientists and ML engineers

creates stickiness for the whole

pipeline. We’re offering a

consistent environment with no

compute infrastructure to manage.

There is high compute upside in

getting these remote, executable

jobs (ML Jobs) that occur any time
data needs to be preprocessed
(new data comes in to be scored)
on container runtime.

How do you automate the model
preprocessing pipeline? What
separate tools are you using for
that orchestration?

How much of your engineering
team's time is spent maintaining
them versus building new ones?

How do you currently manage and
scale the compute resources for
your automated ML jobs?

© 2025 Snowflake Inc. All Rights ReservedDevelop

& Iterate

Orchestrate

& Automate

Manage

Deploy

& Serve

Monitor

What is it? Why it Matters Discovery Questions

This is a centralized way to

discover, manage, and govern all ML assets related to a model, regardless of where it was created.

Models are static in that the training data, the hyperparameters used, the model type, are all point in time.

Model versions and their assets for a particular experiment or model run can vary greatly.

Connecting models to business value is becoming more important and less black box

Snowflake Model Registry and ML Lineage are uniquely positioned to add immediate value.

Low compute upside - more about mindshare of where models are managed

How do you manage models that are in production? Is it in a centralized location or spread

across teams & git repositories?

How do you trace lineage back to

the exact features and data

version used to train it?

How do you track model versions

and who authorizes that?

Could a new teammate find,

understand, and manage model

versions?

© 2025 Snowflake Inc. All Rights ReservedDevelop

& Iterate

Orchestrate

& Automate

Manage

Deploy

& Serve

Monitor

What is it? Why it Matters Discovery Questions

This is essentially the inference

step. It means moving a model

version into production to make

predictions on new data coming

in.

This model can be “served” to

accept new predictions in different ways - usually either in batch (make predictions overnight for the day's data) or real time.

This is where all the magic happens in terms of translating AI into business value.

There is medium compute upside in owning where the inference occurs, what happens with the predictions themselves downstream, and the feature engineering pipeline.

Snowflake has two model serving options - warehouse for batch and SPCS for low latency. We also have a built-in feature store.

Do you currently have models in production today and what percentage are batch vs. near real time?

How do you ensure feature processing pipelines for new data you are scoring?

How do you manage the
underlying compute resources for
your deployed models (scale)?

How do you generate datasets for
training / inference?

© 2025 Snowflake Inc. All Rights ReservedDevelop

& Iterate

Orchestrate

& Automate

Manage Deploy

& Serve

Monitor

What is it? Why it Matters Discovery Questions

Once a model is in production and
making predictions, you must
monitor to ensure it continues to
perform well.

Performing well generally means
that the chosen accuracy metric
remains consistent. Changes in
business processes or underlying
data (drift) can cause model
degradation.

The success of the use case

depends on the ability to
successfully monitor models in
production.

There is low compute upside in
this phase, customer success in
relation to business value is
pertinent though.

Snowflake now has ML

Observability & ML

Explainability to round out a
complete ML lifecycle.

© 2025 Snowflake Inc. All Rights Reserved

How do you track model
performance over time and what
actions are triggered?

How do you detect and get alerted
to model or data drift?

When a model's performance
degrades, what is your process for
troubleshooting the issue and
triggering a retrain?

What is your process for auditing
your ML models for compliance
and explaining why a certain

prediction was made? Compete: AWS

Strengths ♦♦♦♦♦

- Sagemaker ML Platform is “mature” in both capabilities and product marketing
- No gaps in terms of their ML features or tools
- Their ecosystem is sticky and early Snowflake AI story makes AWS seem like our big brother.

© 2022 Snowflake Inc. All Rights Reserved

Weaknesses

How can ❄ Win?

- It really isn't a single platform - while all the services exist, you still have to stitch them together yourself
- Emphasize architecture simplicity
- Reduce time to value by starting and ending in Snowflake
- Architecture complexity to get a model off the ground
- Forced data movement &

egress

- Land & Expand in the ML pipeline

○ If predictions come back to

Snowflake, why not use

our model registry for

inference (on externally

trained models)? Compete: Databricks

Strengths How can ❄ Win?

Weaknesses



- “Industry Leading” in ML

mindshare

- Overwhelming and complex

depending on customer’s ML

maturity

- Grab their attention - find an

opening to showcase specific

capabilities

- Robust ML Ops framework

/ experimentation

- Spark experience required to

build pipelines, cluster

optimization

- Snowflake ML jobs for remote code

execution (chained notebooks
concept)

- Made for ML practitioners

(lots of horsepower)

- Lack of cost transparency to
track ML projects from start to
finish

- Position multi-modal offerings - AI

SQL capabilities combined with

Snowflake ML

© 2022 Snowflake Inc. All Rights ReservedCompete: Microsoft

Strengths Weaknesses



- Really good for customers

who are all in on the

Microsoft ecosystem

- They are effectively single cloud

when it comes to Azure ML

- Tightly integrated to all

Microsoft services and/or

databricks

- For customers leveraging both

Databricks ML and Azure ML, the

line in the sand on separation of

duties hurts their overall

messaging

- Mature ML platform and

great for practitioners

- If using delta lake this can add

complexity for teams and

governance challenges

© 2022 Snowflake Inc. All Rights Reserved

How can * Win?

- If delta lake users and proud of

“open source”, challenge them

to look at their Azure centric

stack

- Pitch multi-cloud

- Find pain points / governance

challenges in bronze layer

- Land & Expand playKey Takeaways & Call to Action

© 2022 Snowflake Inc. All Rights ReservedTHANK YOU

© 2022 Snowflake Inc. All Rights Reserved